

Soar Workshop

RL Tutorial

May 6, 2019

1. RL as a learning mechanism
2. Architecture & agent design
3. Eater integration

Broader Learning Context

From Yann LeCun

■ “Pure” Reinforcement Learning (cherry)

- ▶ The machine predicts a scalar reward given once in a while.
- ▶ **A few bits for some samples**

■ Supervised Learning (icing)

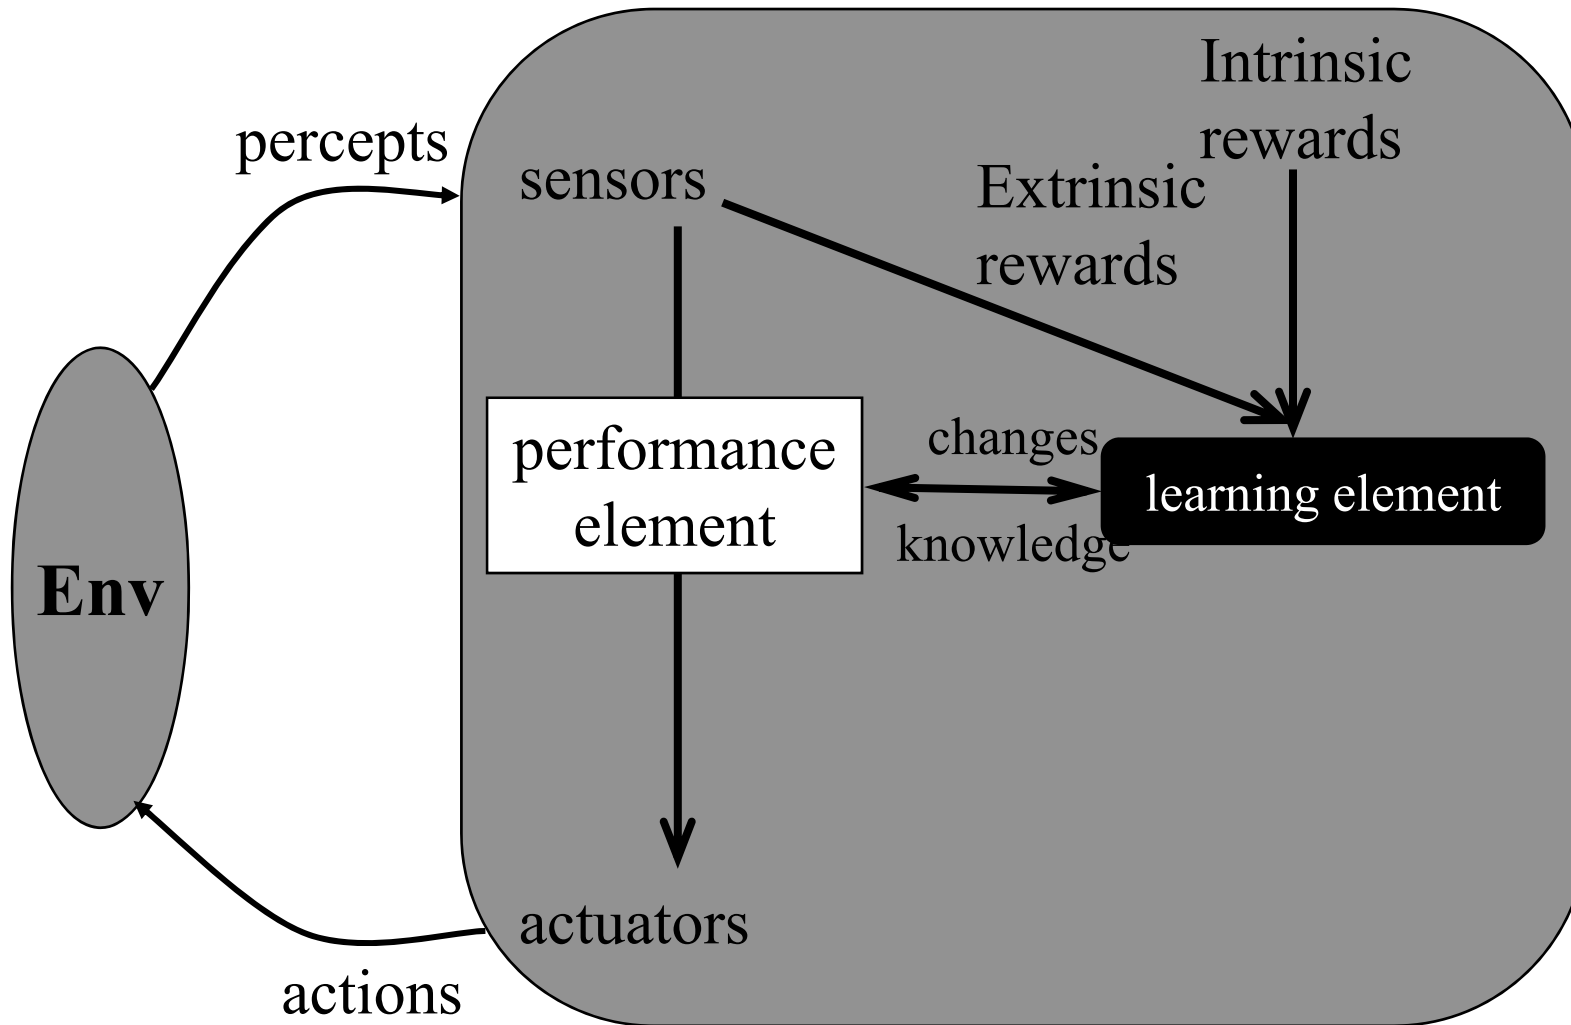
- ▶ The machine predicts a category or a few numbers for each input
- ▶ Predicting human-supplied data
- ▶ **10→10,000 bits per sample**

■ Unsupervised/Predictive Learning (cake)

- ▶ The machine predicts any part of its input for any observed part.
- ▶ Predicts future frames in videos
- ▶ **Millions of bits per sample**



Reinforcement Learning Agent



Reinforcement Learning Agent

- Learns by doing!
- How animals learn by (lots of) training
- Good for optimizing decision making
- Requires frequent reward
- RL makes explicit the *trade-off* between
 - Exploration: acting to *learn* the environment,
 - Exploitation: acting to *maximize reward*.

What Has RL Done?

- Game playing
 - A world-class backgammon player
 - A checkers player in the 1950s
 - Lots of video games in 2015-2016
 - Alpha-Go, Alpha-Go Zero, Alpha-Star in 2016-2019
- Robotics
 - A faster walking algorithm for the Aibo
 - A stable controller for a passive-dynamic walker

Mazes: Demo




What is Reinforcement Learning (RL)?

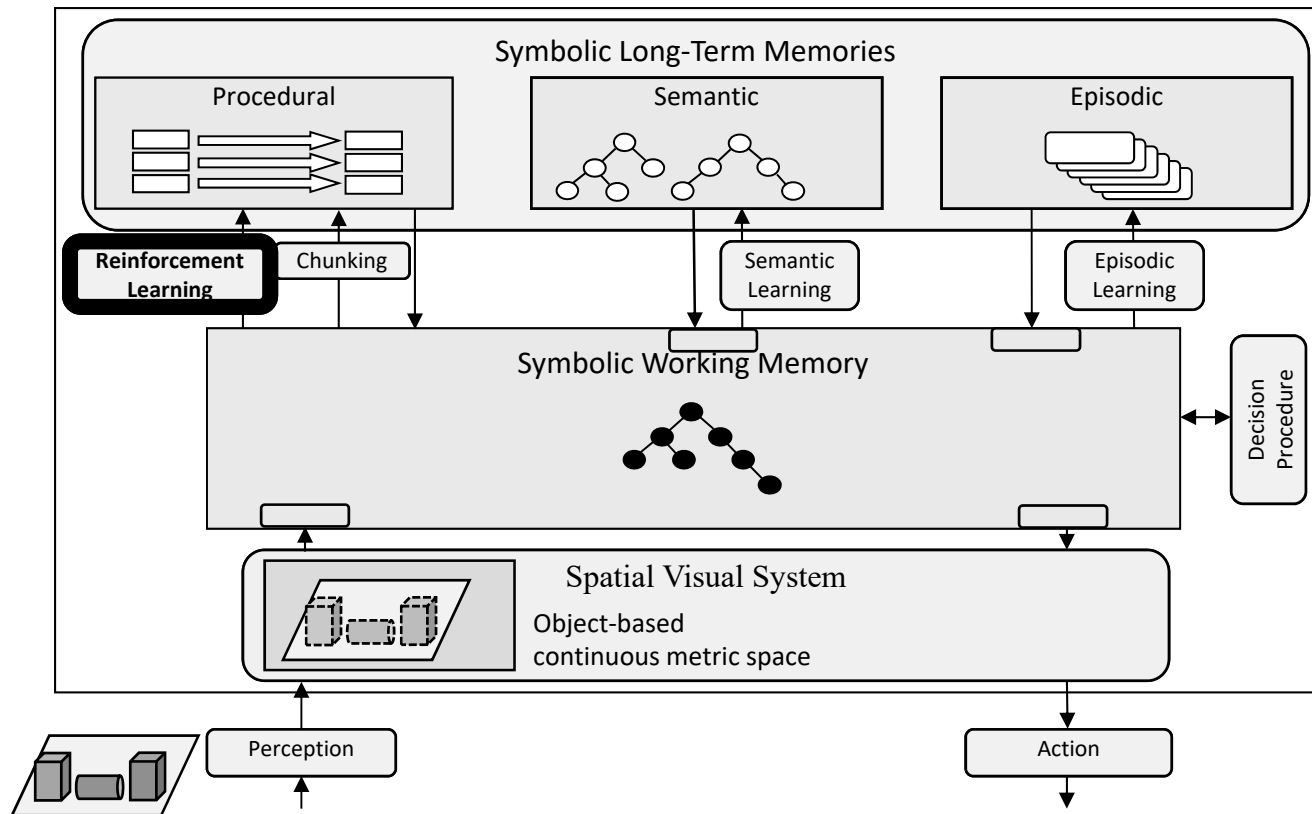
- Goal: learn an optimal action **policy**; given an environment that provides states, affords actions, and provides feedback as numerical **reward**, maximize the expected future reward.
- Typically involves learning a **value function** that maps states (or state-action pairs) to a prediction of expected future reward.
- Allows giving reward for achieving goal and having system figure out how to achieve reward.
- In Soar, RL involves learning operator *selection* knowledge: numeric preferences.

RL Cycle

Goal: learn an action-selection policy such as to maximize expected receipt of future reward




Soar 9



Methods for Learning Procedural Knowledge


Chunking

- Converts *deliberation* in substates into *reaction* via rule compilation



Reinforcement Learning

- *Tunes* operator numeric preferences to reflect expectation of reward

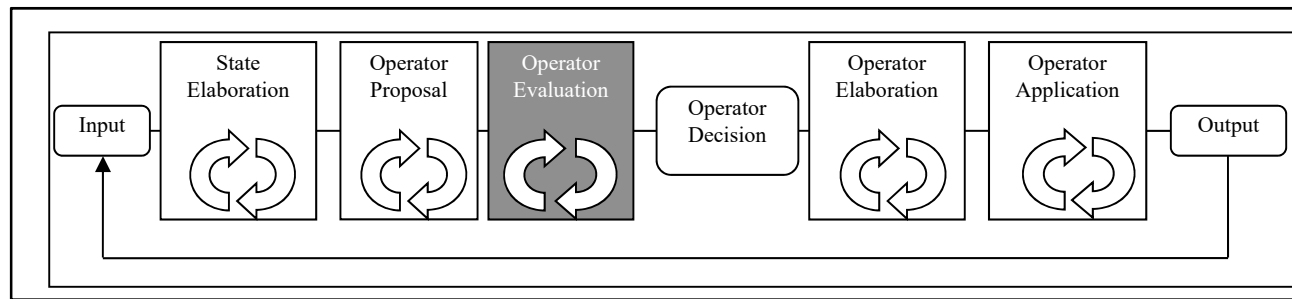


Can be used together
Chunking can learn initial RL rules

- Creates new rules
- Updates existing rules

Soar Basic Functions

- 1. Input from environment
2. Elaborate current situation: *parallel rules*
3. Propose operators via acceptable preferences
4. Evaluate operators via *preferences: Numeric indifferent preference*
5. Select operator
6. Apply operator: Modify internal data structures: *parallel rules*
7. Output to motor system [and access to long-term memories]



Left-Right Demo

1. Soar Java Debugger
2. Source `left-right.soar` file



Left-Right Demo

Script

1. `srand 50412`
2. `step`
3. `step`
4. `click: op_pref tab`
 - note numeric indifferent preferences
5. `print left-right*rl*left`
6. `print left-right*rl*right`
7. `step`
 - note movement direction
8. `print left-right*rl*left`
9. `print left-right*rl*right`
10. `init-soar`
11. Repeat from #2 (~5 times)

Left-Right: Takeaways

Reinforcement learning changes rules in procedural memory

- Changes are persistent (until changed by new update).
- Change affects *numeric indifferent preferences*, which in turn affects the selection of operators.
- Change is in the direction of the underlying reward signal (will discuss this more shortly)

RL -> Architecture & Agent Design

Value function

*via RL rules **[agent]***

Reward

*via working-memory structures **[architecture, agent]***

Policy updates

*via Temporal Difference (TD) Learning **[architecture]***

RL Rules

The RL mechanism maintains Q-values for state-operator pairs in operator selection rules, identified by syntax

- Action has a single action, that is a single numeric indifferent preference with a constant value

```
sp {left-right*rl*left
    (state <s> ^name left-right
        ^operator <op> +)
    (<op> ^name move
        ^dir left)
-->
(<s> ^operator <op> = 0) }
```

```
sp {left-right*rl*right
    (state <s> ^name left-right
        ^operator <op> +)
    (<op> ^name move
        ^dir right)
-->
(<s> ^operator <op> = 0) }
```


Left-Right Demo

Focus: RL Rules

1. Soar Java Debugger
2. **Source** `left-right.soar` **file**
3. `print --full --rl`
4. `run`
5. `print --full --rl`
6. `print --rl`

Reward Representation


Each state in working memory has a `reward-link` structure

Reward is recognized by syntax


```
(<s> ^reward-link <r-link>)  
(<r-link> ^reward <r>)  
(<r> ^value [integer or float])
```


- The reward-link is **not** directly modified by the environment or architecture
- Reward is collected at the beginning of each *decide* phase
- Reward on a state's reward-link pertains only to that state (more on this later)
- Reward can come from multiple rules: reward values are summed by default

Reward Rule Examples


```
sp {left-right*reward*left
  (state <s> ^name left-right
    ^location 
    ^reward-link <rl>)
```

-->

```
(<rl> ^reward <r>)
(<r> ^value  ) }
```

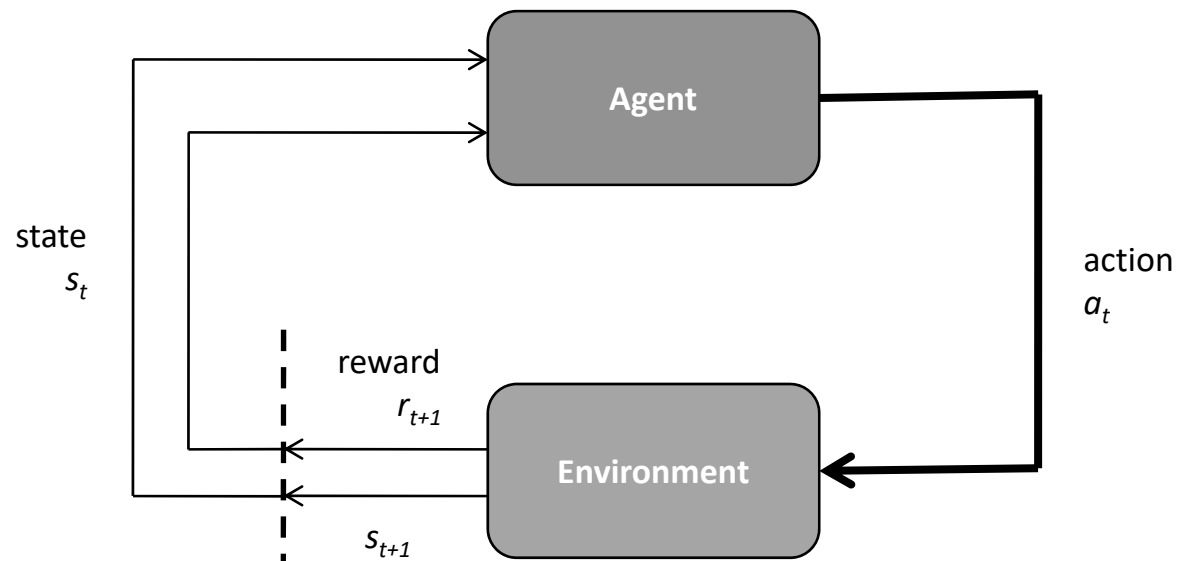
```
sp {left-right*reward*right
  (state <s> ^name left-right
    ^location 
    ^reward-link <rl>)
```

-->

```
(<rl> ^reward <r>)
(<r> ^value  ) }
```



RL Cycle



RL Cycle in Soar

	Input	Propose	Decide	Apply	Output
d					
d+1					

RL Cycle in Soar

	Input	Propose	Decide	Apply	Output
d	state _d				
d+1					

RL Cycle in Soar

	Input	Propose	Decide	Apply	Output
d	$state_d$	evaluate $operators_d$			
d+1					

RL Cycle in Soar

	Input	Propose	Decide	Apply	Output
d	$state_d$	evaluate $operators_d$	select $operator_d$		
d+1					

RL Cycle in Soar

	Input	Propose	Decide	Apply	Output
d	$state_d$	evaluate $operators_d$	select $operator_d$		initiate external action(s)
d+1					

RL Cycle in Soar

	Input	Propose	Decide	Apply	Output
d	$state_d$	evaluate $operators_d$	select $operator_d$		initiate external action(s)
d+1	$state_{d+1}$ $reward_{d+1}$				

RL Cycle in Soar

	Input	Propose	Decide	Apply	Output
d	$state_d$	evaluate operators _d	select operator _d		initiate external action(s)
d+1	$state_{d+1}$ $reward_{d+1}$	evaluate operators _{d+1}			

RL Cycle in Soar

	Input	Propose	Decide	Apply	Output
d	state _d	evaluate operators _d	select operator _d		initiate external action(s)
d+1	state _{d+1} reward _{d+1}	evaluate operators _{d+1}	select operator _{d+1} update policy _d		

RL Updates

- Takes place during *decide* phase, after operator selection
- For all RL rule instantiations (**n**) that supported the *last selected* operator

$$\text{value}_{d+1} = \text{value}_d + (\delta_d / n)$$

Where, roughly...

$$\delta_d = \alpha [\text{reward}_{d+1} + \gamma(q_{d+1}) - \text{value}_d]$$

Where...

- α is a parameter (learning rate)
- γ is a parameter (discount rate)
- q_{d+1} is dictated by learning policy
 - On-policy (SARSA): value of selected operator
 - Off-policy (Q-learning): value of operator with maximum selection probability

Value Function

Issues

Structure

- What features comprise RL-rule conditions (tradeoff: convergence speed vs. performance)
- Lots of features -> computationally infeasible
- Few features -> not specific enough

Initialization

- Quality estimates may bootstrap agent performance and reduce time to convergence
- Set initial values of the RL rules.

Eaters RL

- General idea:
 - RL rules will learn to select between forward and rotate operators.

Eaters RL 1

Get your eater code

Add to top of file or

create a new file (eater-RL.soar)

– turn on RL

- **rl -s learning on**
- **Indifferent-selection -b** # use boltzman decision making

Eaters RL 2

Remove indifferent preference from proposals so RL rules will influence decision.

```
sp {random*propose*forward
    (state <s> ^name eater
        ^io.input-link.front)
-->
    (<s> ^operator <op> +) ←
    (<op> ^name forward)}
```



```
sp {random*propose*rotate
    (state <s> ^name eater
        ^io.input-link.front)
-->
    (<s> ^operator <op> +) ←
    (<op> ^name rotate)}
```

Just add these to a new file and they will load over your old versions.

Eaters RL 3

Generate RL rules for every color and operator combination:

```
gp {eater*evaluate*forward
  (state <s> ^name eater
    ^io.input-link.front [ red wall blue empty green purple ]
    ^operator <op1> +)
  (<op1> ^name forward)
-->
  (<s> ^operator <op1> = 0.0) }
```

```
gp {eater*evaluate*rotate
  (state <s> ^name eater
    ^io.input-link.front [ red wall blue empty green purple ]
    ^operator <op1> +)
  (<op1> ^name rotate)
-->
  (<s> ^operator <op1> = 0.0) }
```

Each of these will generate 6 rules!

RL will change the value of = 0.0 in each of the rules as it learns

Eaters RL 4

Add rule that assigns reward – use the change in score:

```
sp {eater*elaborate*state
    (state <s> ^name eater
        ^reward-link <rl>
        ^io.input-link.score-diff <d>)
-->
    (<rl> ^reward.value <d>)
}
```

Run!

- Run eater
- Look at rl rules: $p - r$
- Reset eater (type “r”), run again
- See how rl rules change:
 - Number of updates
 - Value of indifferent preference
- Gets better, but is very limited by the operators available (forward and rotate).